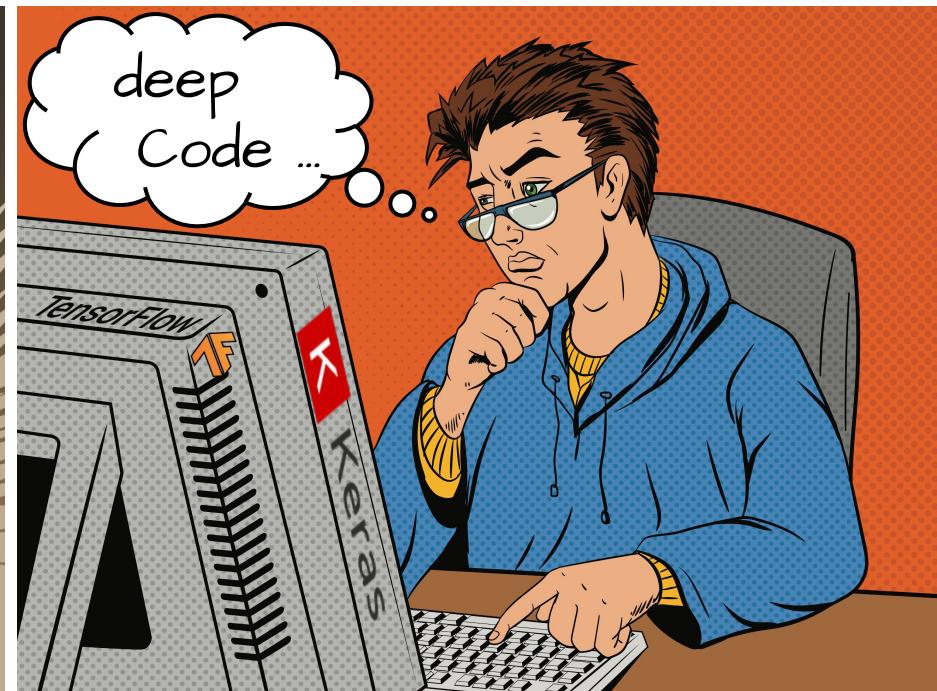
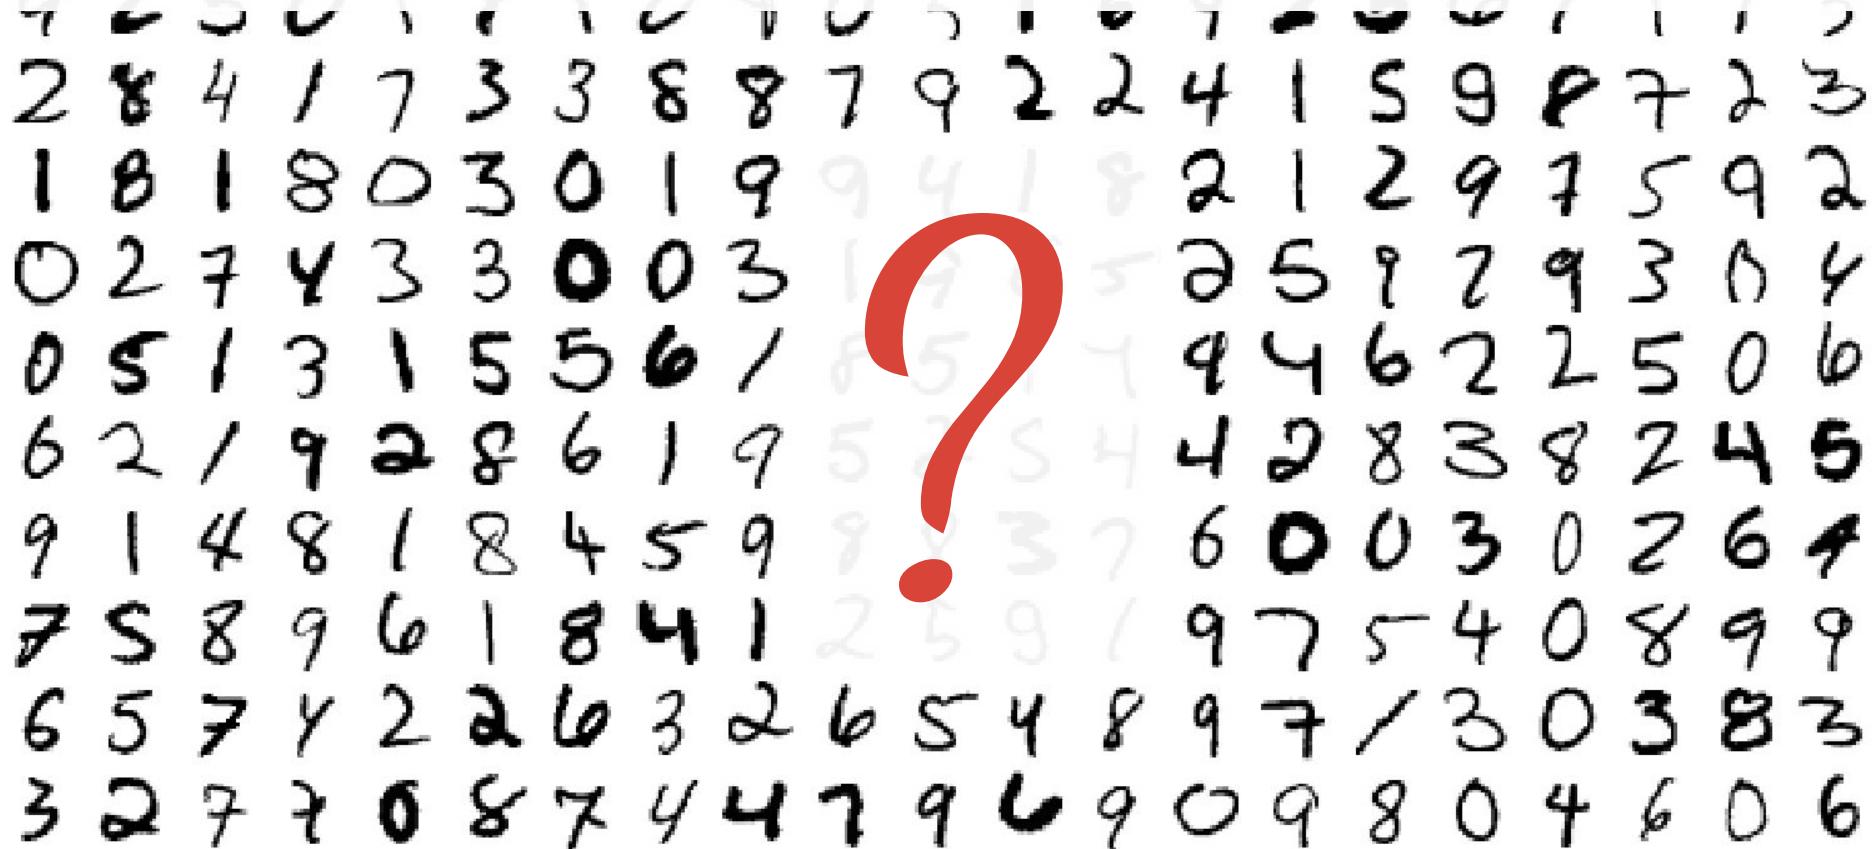


>TensorFlow, Keras and deep learning_

without a PhD

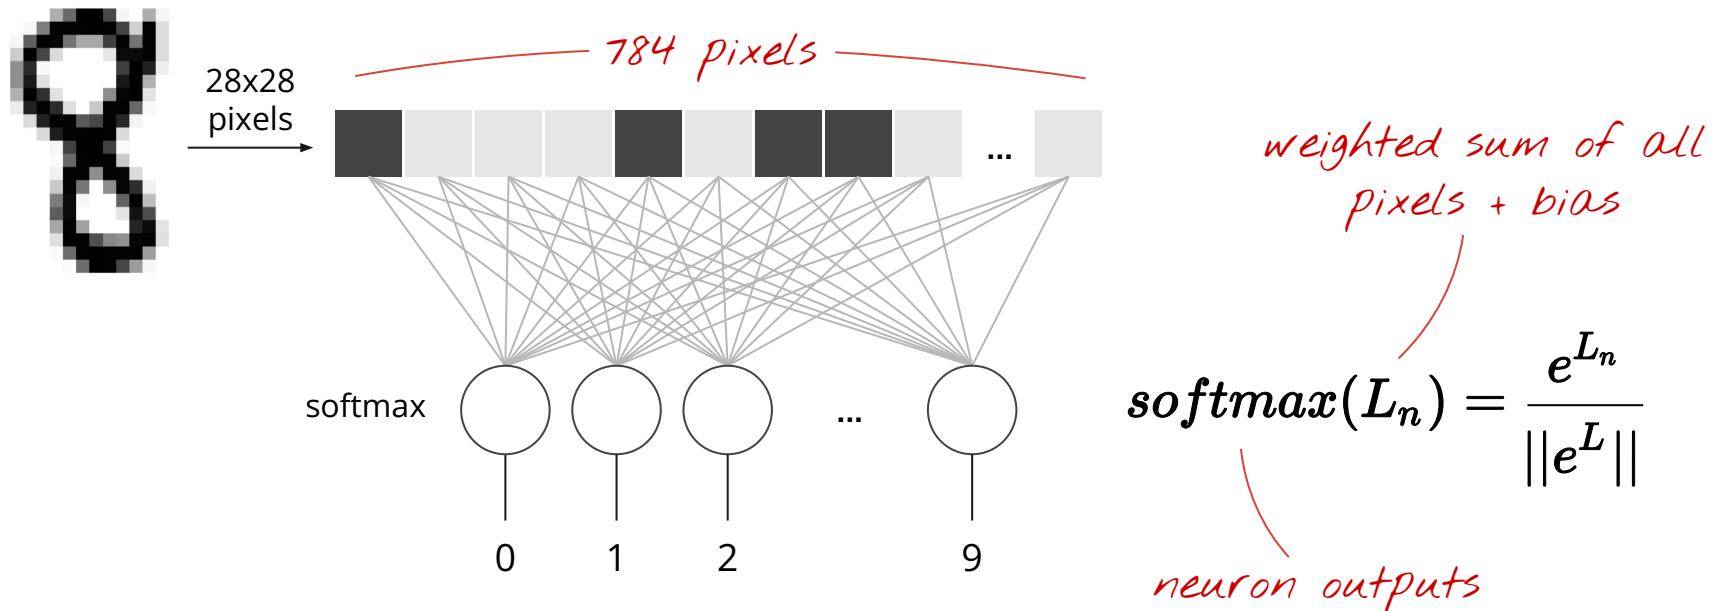


Hello World: handwritten digits classification - MNIST

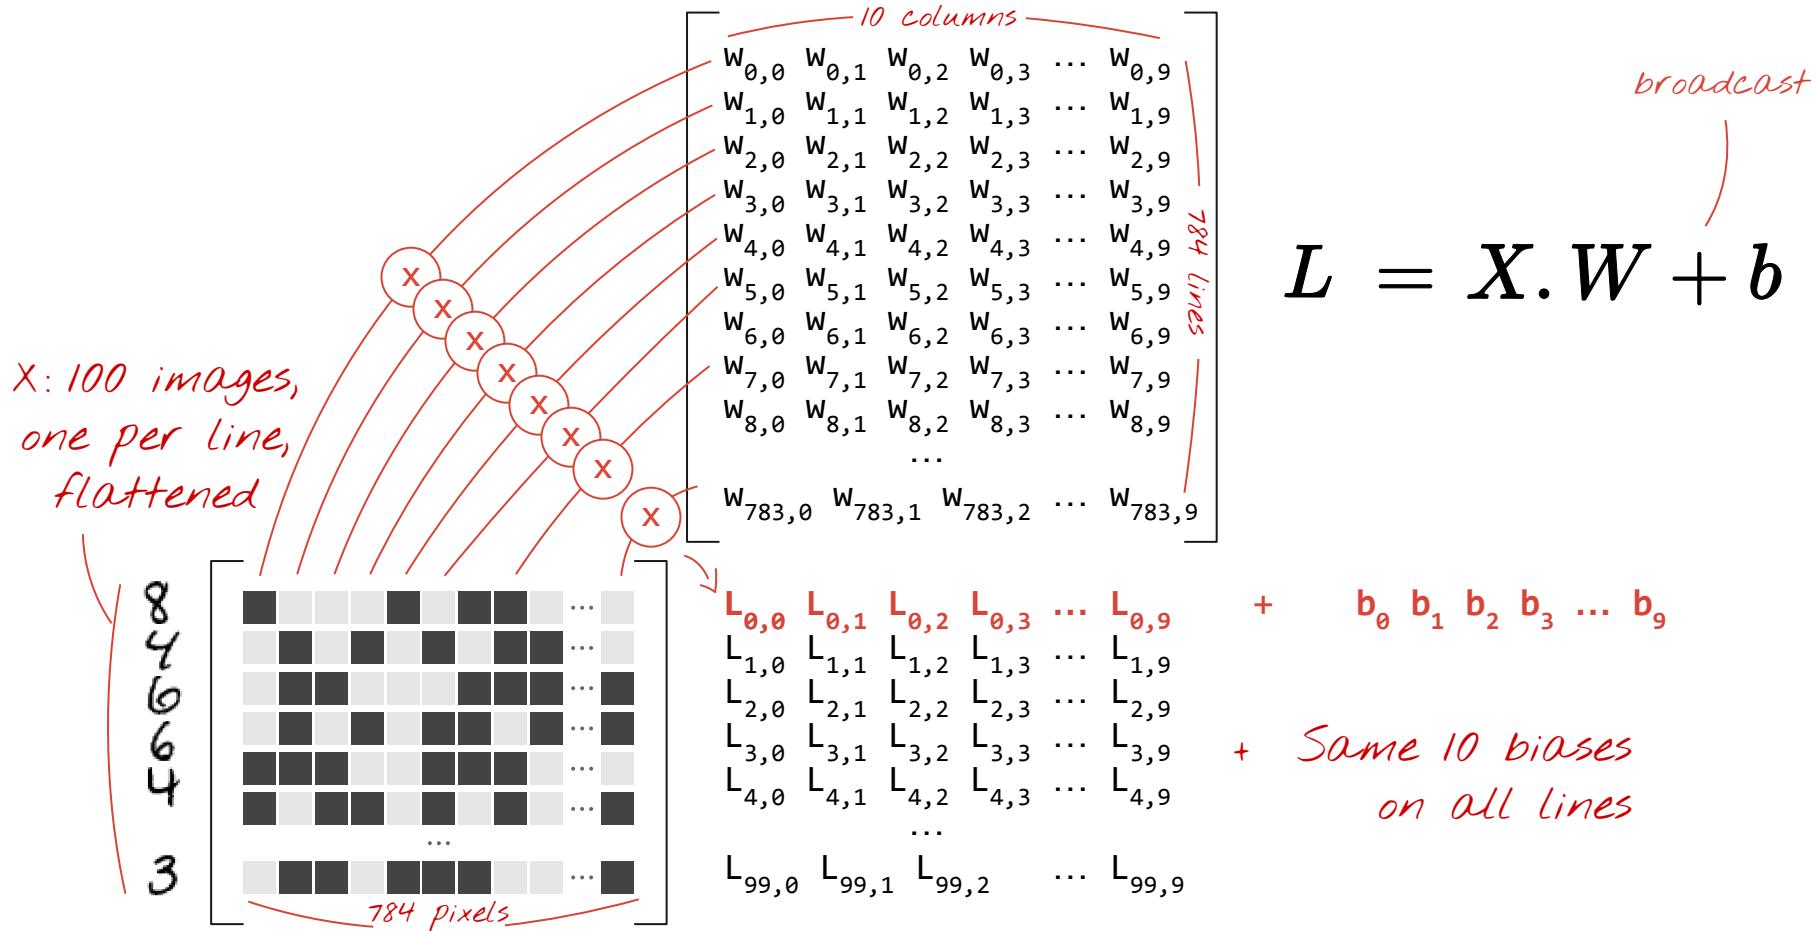


MNIST = Mixed National Institute of Standards and Technology - Download the dataset at <http://yann.lecun.com/exdb/mnist/>

Very simple model: softmax classification



In matrix notation, 100 images at a time



Softmax, on a batch of images

Predictions

$Y[100, 10]$

$$Y = \text{softmax}(X \cdot W + b)$$

applied line
by line

tensor shapes in []

Images

$X[100, 784]$

Weights

$W[784, 10]$

Biases

$b[10]$

matrix multiply

broadcast
on all lines

Now in TensorFlow (Python)

tensor shapes: $X[100, 784]$ $W[784, 10]$ $b[10]$

$Y = \text{tf.nn.softmax}(\text{tf.matmul}(X, W) + b)$

matrix multiply *broadcast
on all lines*

Success ?

0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	1	0	0	0

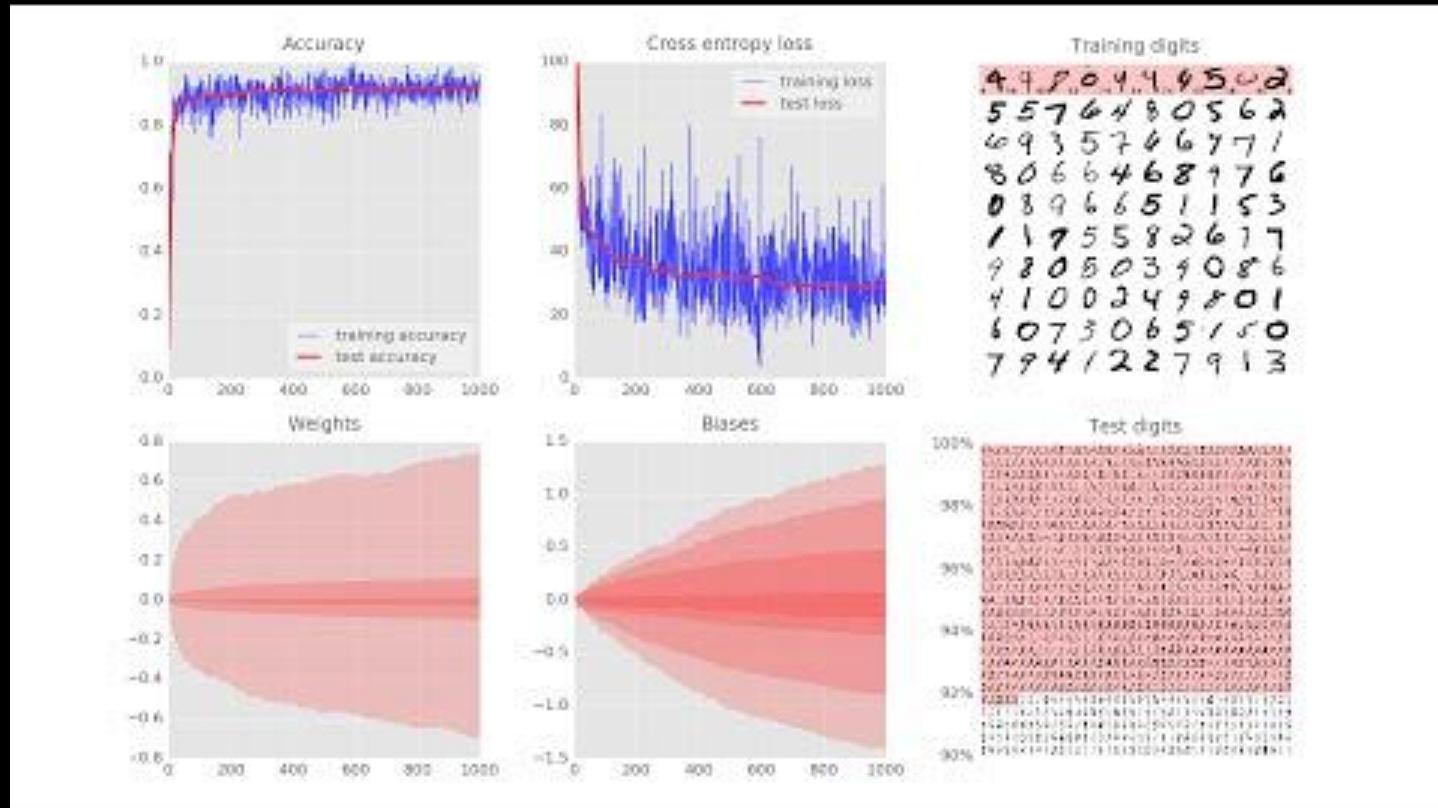
actual probabilities, "one-hot" encoded

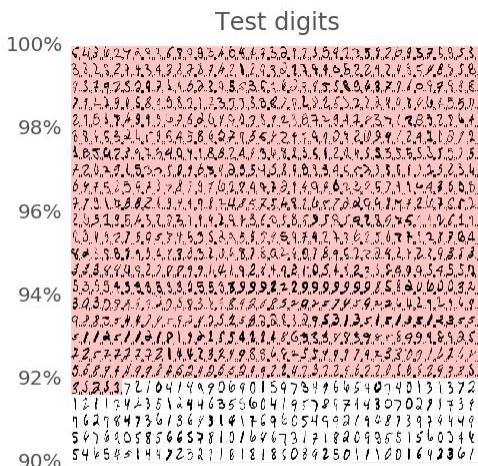
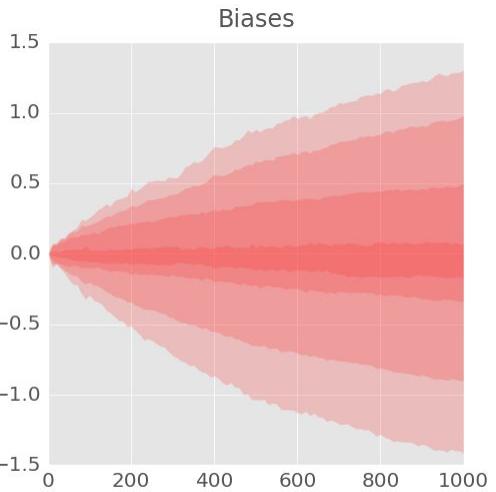
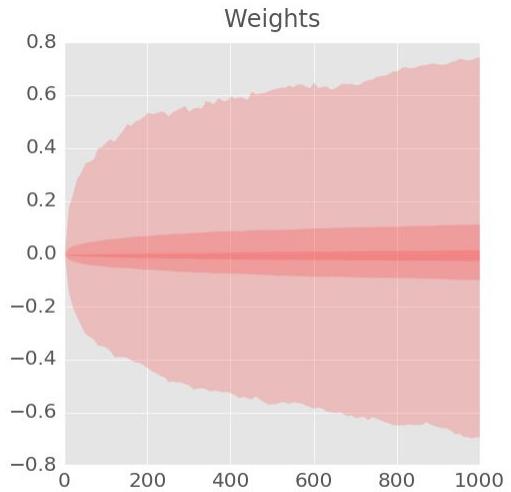
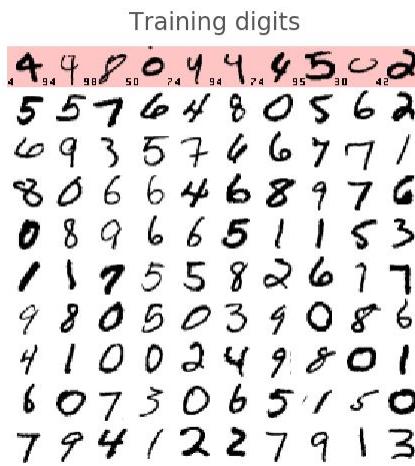
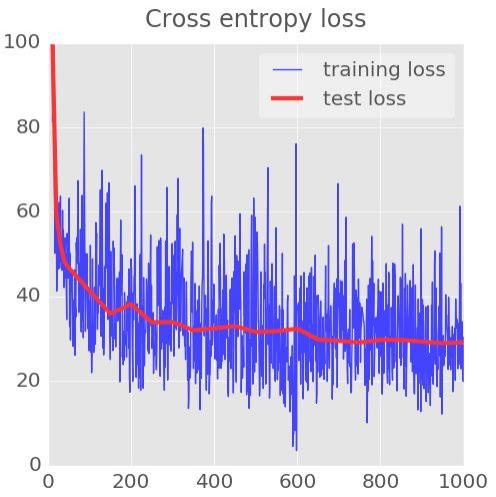
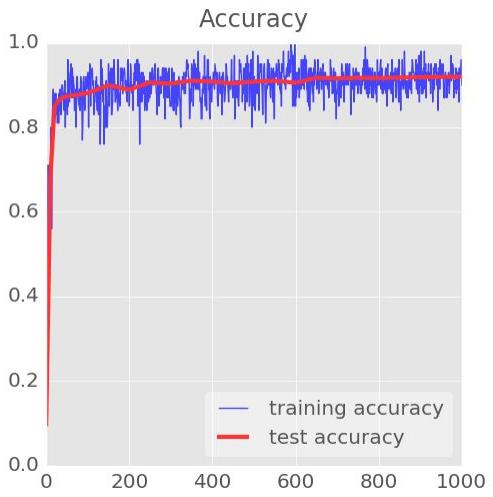
Cross entropy: $-\sum Y'_i \cdot \log(Y_i)$

)
this is a "6"
computed probabilities

.01	.03	.00	.04	.03	.05	0.8	.02	.01	.01
0	1	2	3	4	5	6	7	8	9

Demo





92%

TensorFlow - initialisation

```
import tensorflow as tf  
  
X = tf.placeholder(tf.float32, [None, 28, 28, 1])  
W = tf.Variable(tf.zeros([784, 10]))  
b = tf.Variable(tf.zeros([10]))  
  
init = tf.initialize_all_variables()
```

this will become the batch size, 100
|
28 x 28 grayscale images

Training = computing variables W and b

TensorFlow - success metrics

```
# model  
Y = tf.nn.softmax(tf.matmul(tf.reshape(X, [-1, 784]), W) + b)  
# placeholder for correct answers  
Y_ = tf.placeholder(tf.float32, [None, 10])  
# loss function  
cross_entropy = -tf.reduce_sum(Y_ * tf.log(Y))  
# % of correct answers found in batch  
is_correct = tf.equal(tf.argmax(Y, 1), tf.argmax(Y_, 1))  
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

flattening images

/

"one-hot" encoded

"one-hot" decoding

/

TensorFlow - training

```
optimizer = tf.train.GradientDescentOptimizer(0.003)
train_step = optimizer.minimize(cross_entropy)
```

learning rate



loss function



TensorFlow - run !

```
sess = tf.Session()  
sess.run(init)  
  
for i in range(1000):  
    # Load batch of images and correct answers  
    batch_X, batch_Y = mnist.train.next_batch(100)  
    train_data={X: batch_X, Y_: batch_Y}  
  
    # train  
    sess.run(train_step, feed_dict=train_data)
```

running a Tensorflow computation, feeding placeholders

```
# success ?  
a,c = sess.run([accuracy, cross_entropy], feed_dict=train_data)  
  
# success on test data ?  
test_data={X: mnist.test.images, Y_: mnist.test.labels}  
a,c = sess.run([accuracy, cross_entropy], feed=test_data)
```

Tip:

do this
every 100
iterations

TensorFlow - full python code

```
import tensorflow as tf

X = tf.placeholder(tf.float32, [None, 28, 28, 1])
W = tf.Variable(tf.zeros([784, 10]))
b = tf.Variable(tf.zeros([10]))
init = tf.initialize_all_variables()

# model
Y = tf.nn.softmax(tf.matmul(tf.reshape(X, [-1, 784]), W) + b)

# placeholder for correct answers
Y_ = tf.placeholder(tf.float32, [None, 10])

# loss function
cross_entropy = -tf.reduce_sum(Y_ * tf.log(Y))

# % of correct answers found in batch
is_correct = tf.equal(tf.argmax(Y, 1), tf.argmax(Y_, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

initialisation

model

success metrics

```
optimizer = tf.train.GradientDescentOptimizer(0.003)
train_step = optimizer.minimize(cross_entropy)
```

training step

```
sess = tf.Session()
sess.run(init)
```

```
for i in range(10000):
    # Load batch of images and correct answers
    batch_X, batch_Y = mnist.train.next_batch(100)
    train_data = {X: batch_X, Y_: batch_Y}
```

```
# train
sess.run(train_step, feed_dict=train_data)
```

Run

```
# success ? add code to print it
a, c = sess.run([accuracy, cross_entropy], feed=train_data)
```

```
# success on test data ?
test_data = {X: mnist.test.images, Y_: mnist.test.labels}
a, c = sess.run([accuracy, cross_entropy], feed=test_data)
```

Cookbook

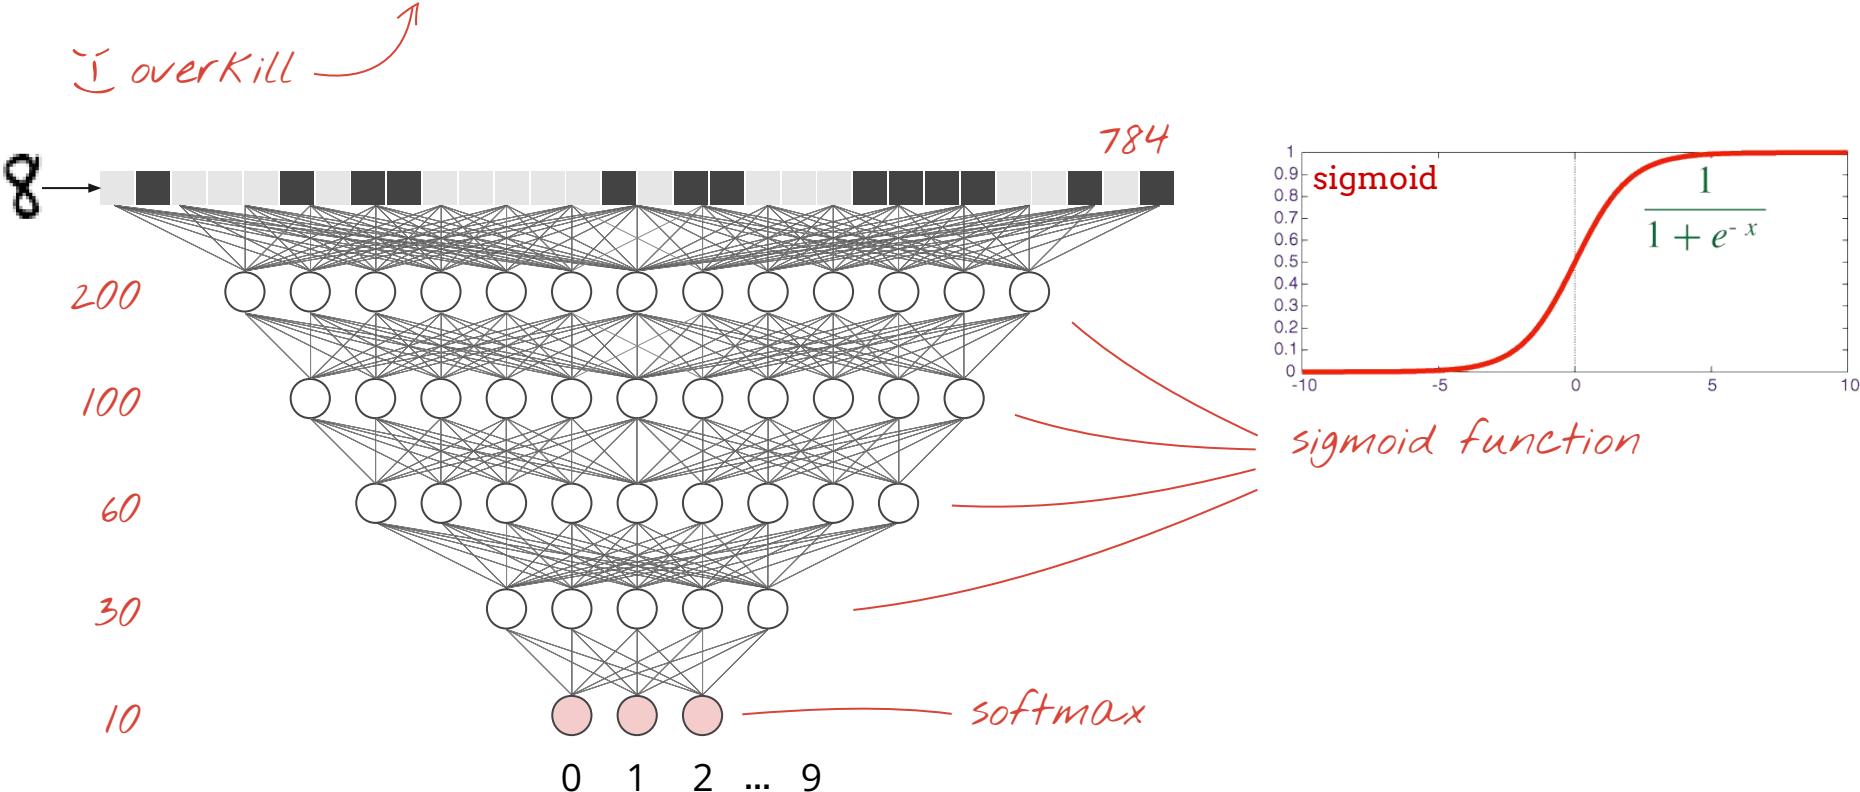
Softmax
Cross-entropy
Mini-batch





Go deep !

Let's try 5 fully-connected layers !



TensorFlow - initialisation

```
K = 200  
L = 100  
M = 60  
N = 30
```

*weights initialised
with random values*

```
W1 = tf.Variable(tf.truncated_normal([28*28, K] ,stddev=0.1))  
B1 = tf.Variable(tf.zeros([K]))
```

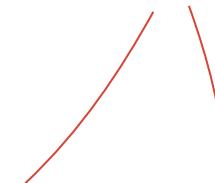
```
W2 = tf.Variable(tf.truncated_normal([K, L], stddev=0.1))  
B2 = tf.Variable(tf.zeros([L]))
```

```
W3 = tf.Variable(tf.truncated_normal([L, M], stddev=0.1))  
B3 = tf.Variable(tf.zeros([M]))  
W4 = tf.Variable(tf.truncated_normal([M, N], stddev=0.1))  
B4 = tf.Variable(tf.zeros([N]))  
W5 = tf.Variable(tf.truncated_normal([N, 10], stddev=0.1))  
B5 = tf.Variable(tf.zeros([10]))
```

TensorFlow - the model

```
X = tf.reshape(X, [-1, 28*28])
```

weights and biases



```
Y1 = tf.nn.sigmoid(tf.matmul(X, W1) + B1)
```

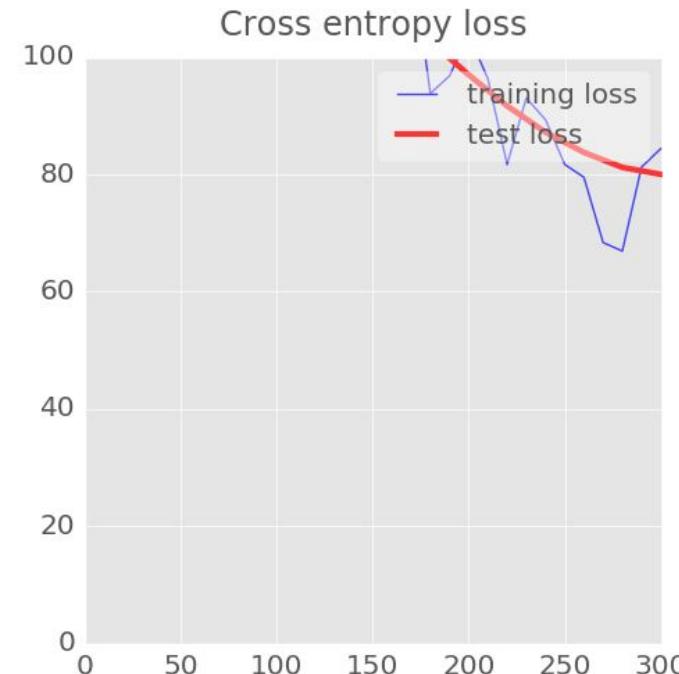
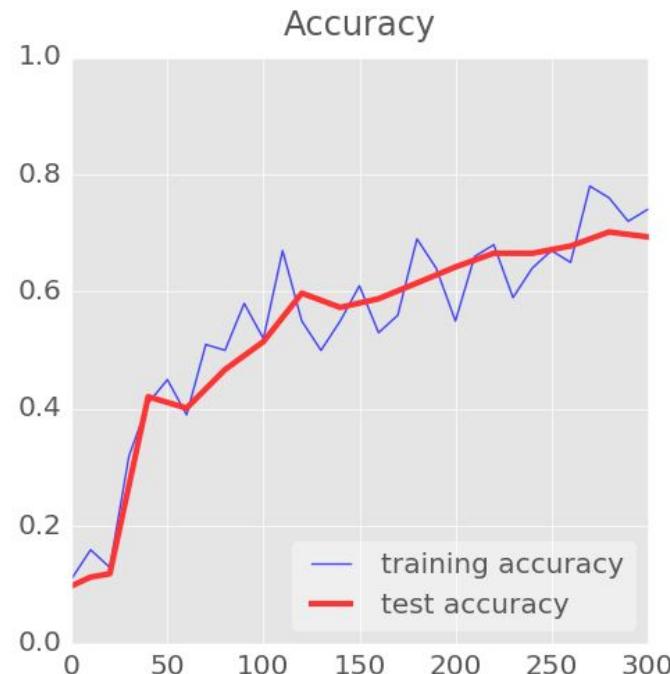
```
Y2 = tf.nn.sigmoid(tf.matmul(Y1, W2) + B2)
```

```
Y3 = tf.nn.sigmoid(tf.matmul(Y2, W3) + B3)
```

```
Y4 = tf.nn.sigmoid(tf.matmul(Y3, W4) + B4)
```

```
Y = tf.nn.softmax(tf.matmul(Y4, W5) + B5)
```

Demo - slow start ?



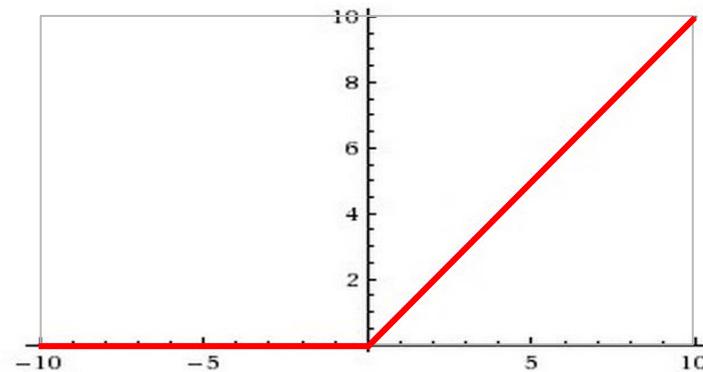
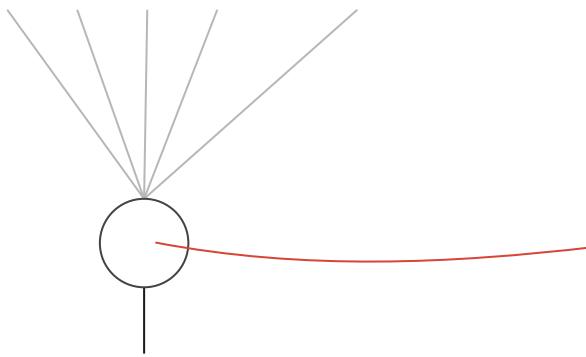


Relu !

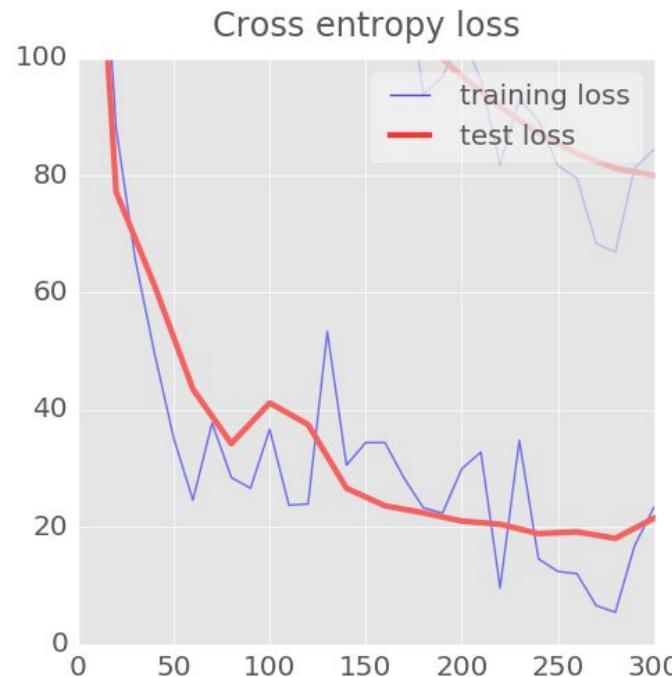
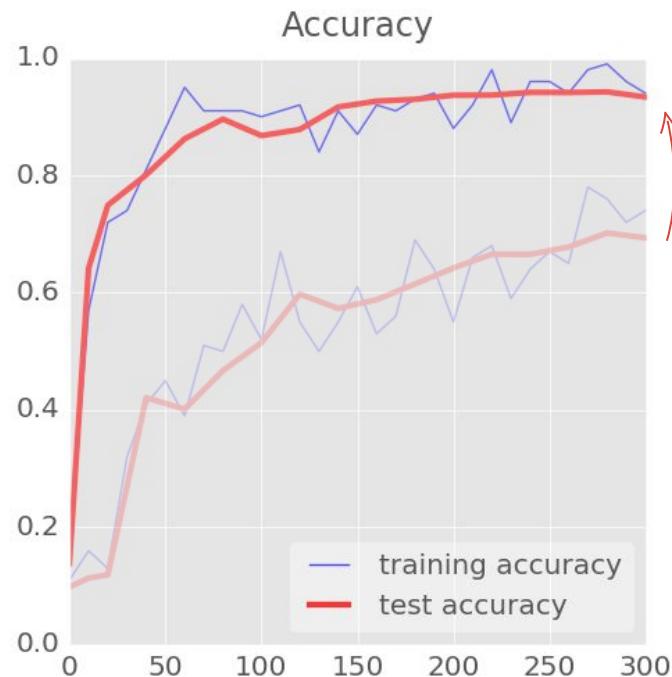


RELU

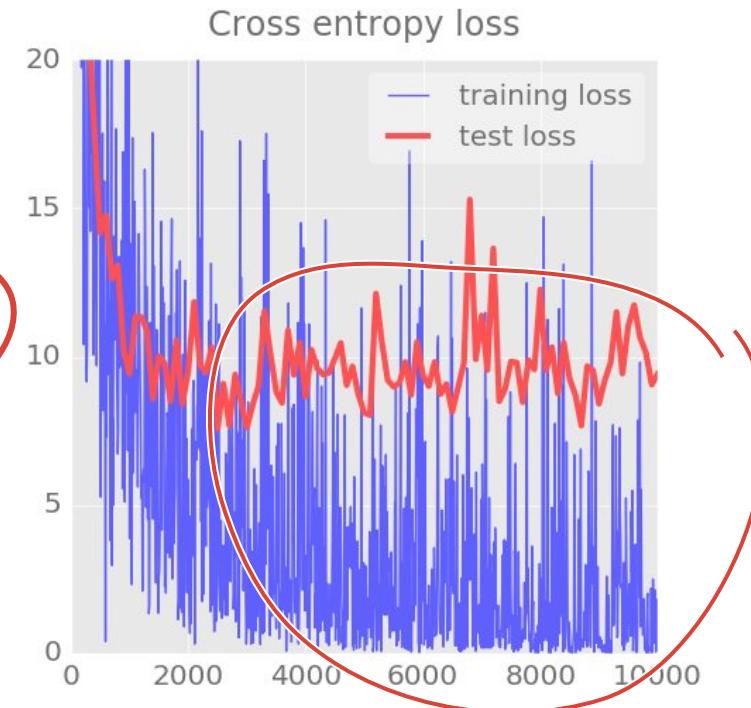
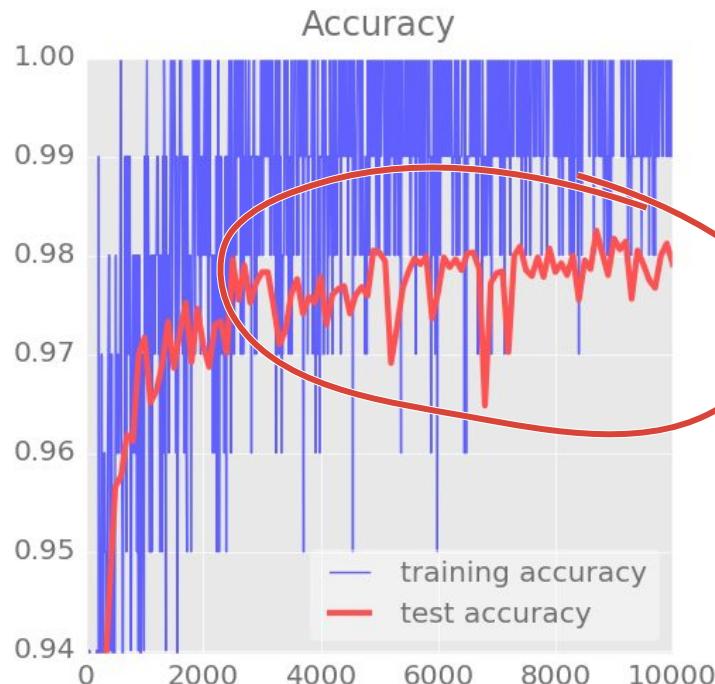
RELU = Rectified Linear Unit


$$Y = \text{tf.nn.relu}(\text{tf.matmul}(X, W) + b)$$

RELU



Demo - noisy accuracy curve ?



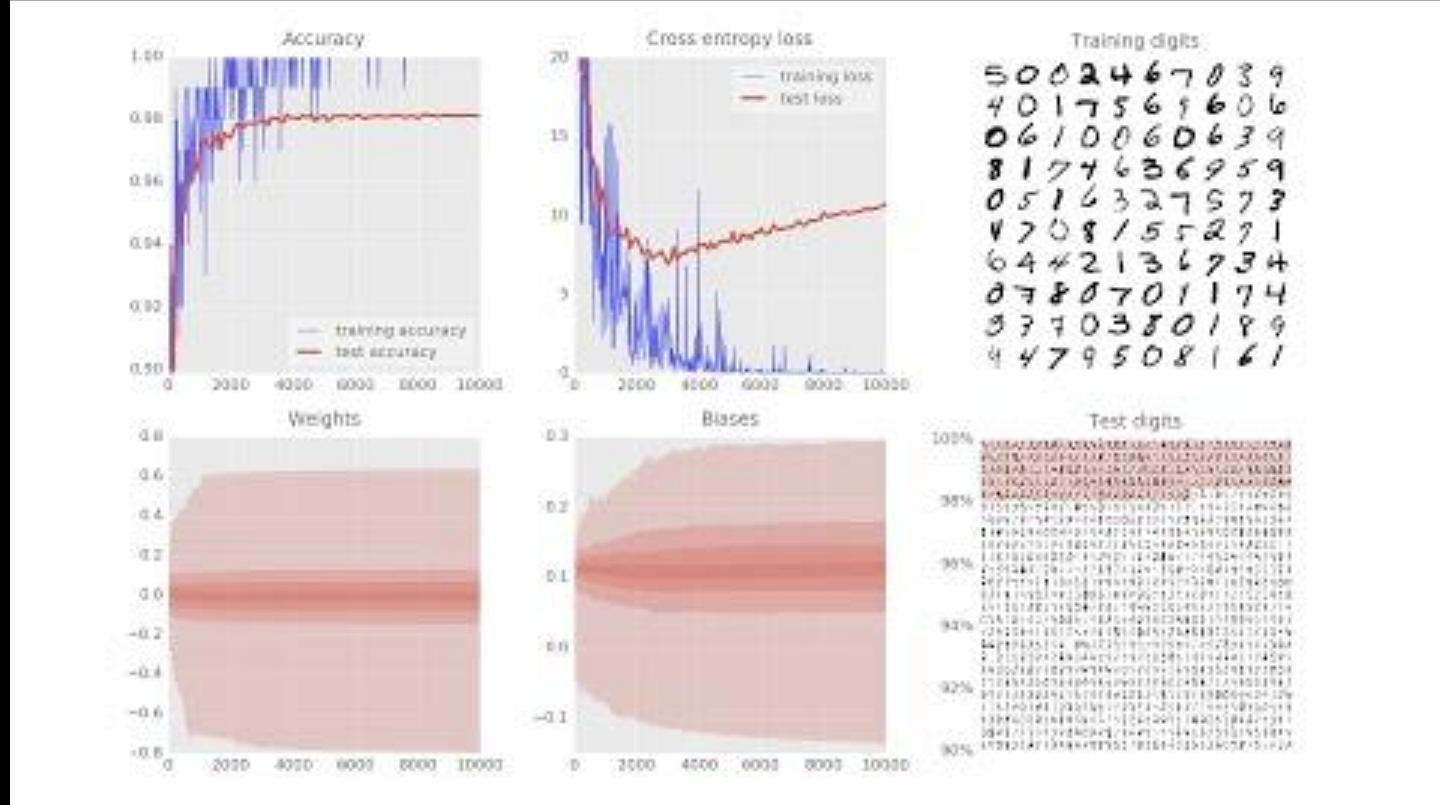
yuck!

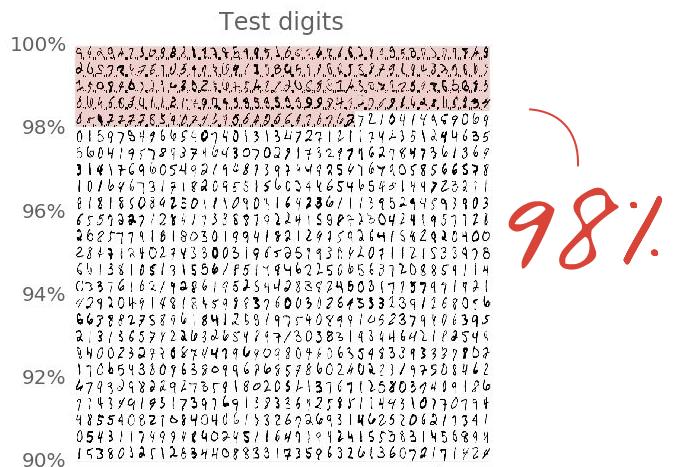
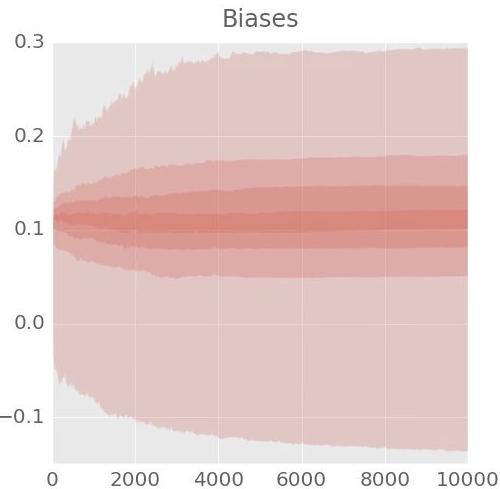
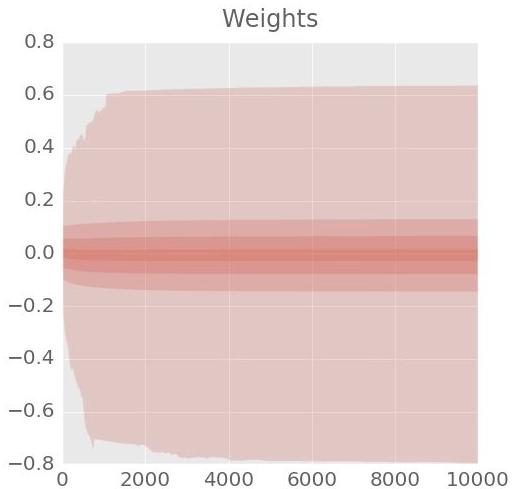
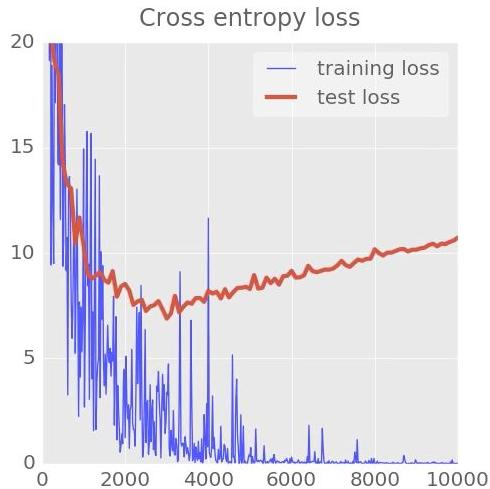
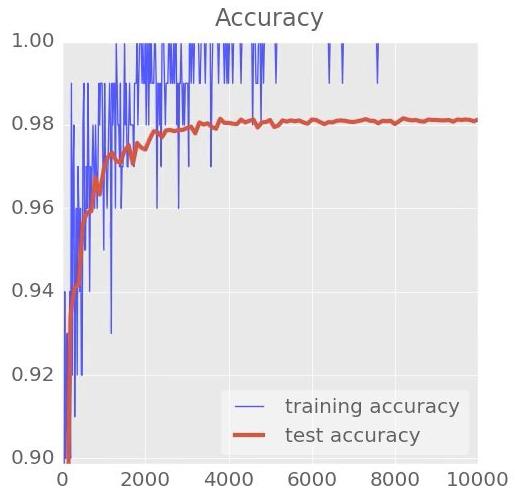
Slow down...

Learning
rate decay

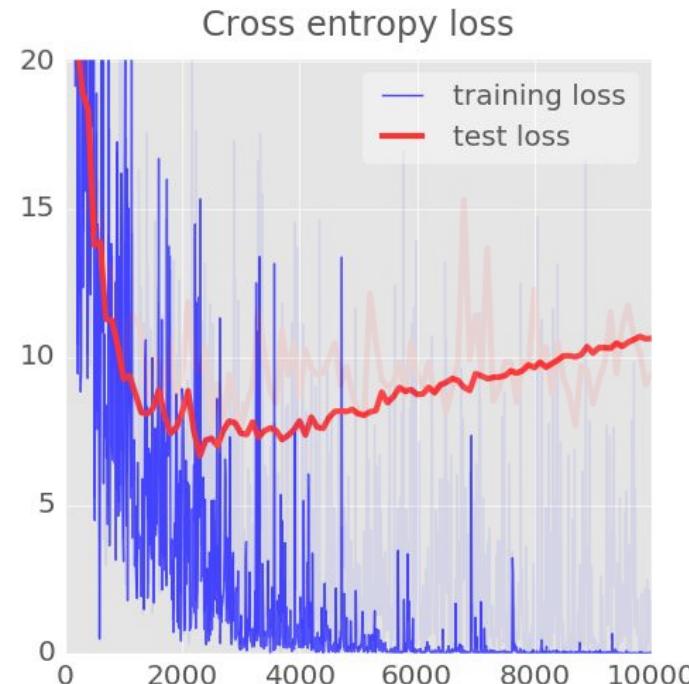
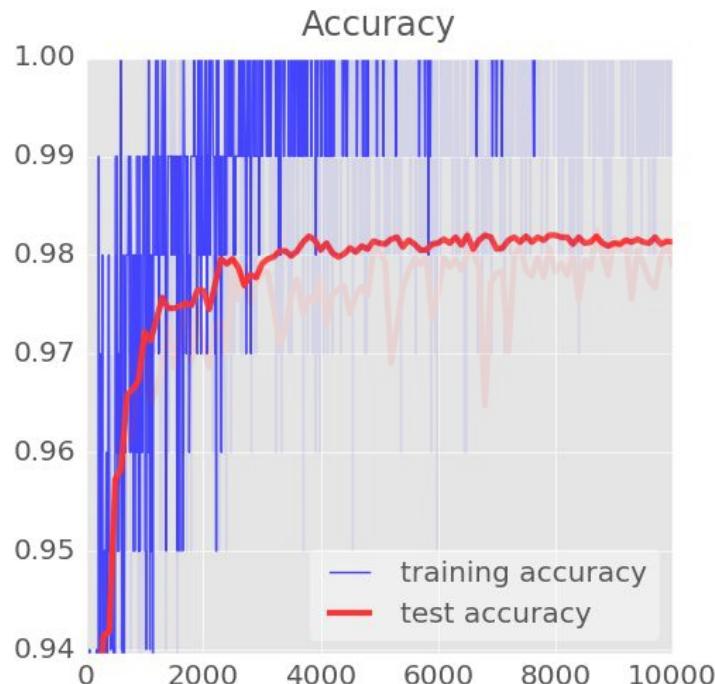


Demo



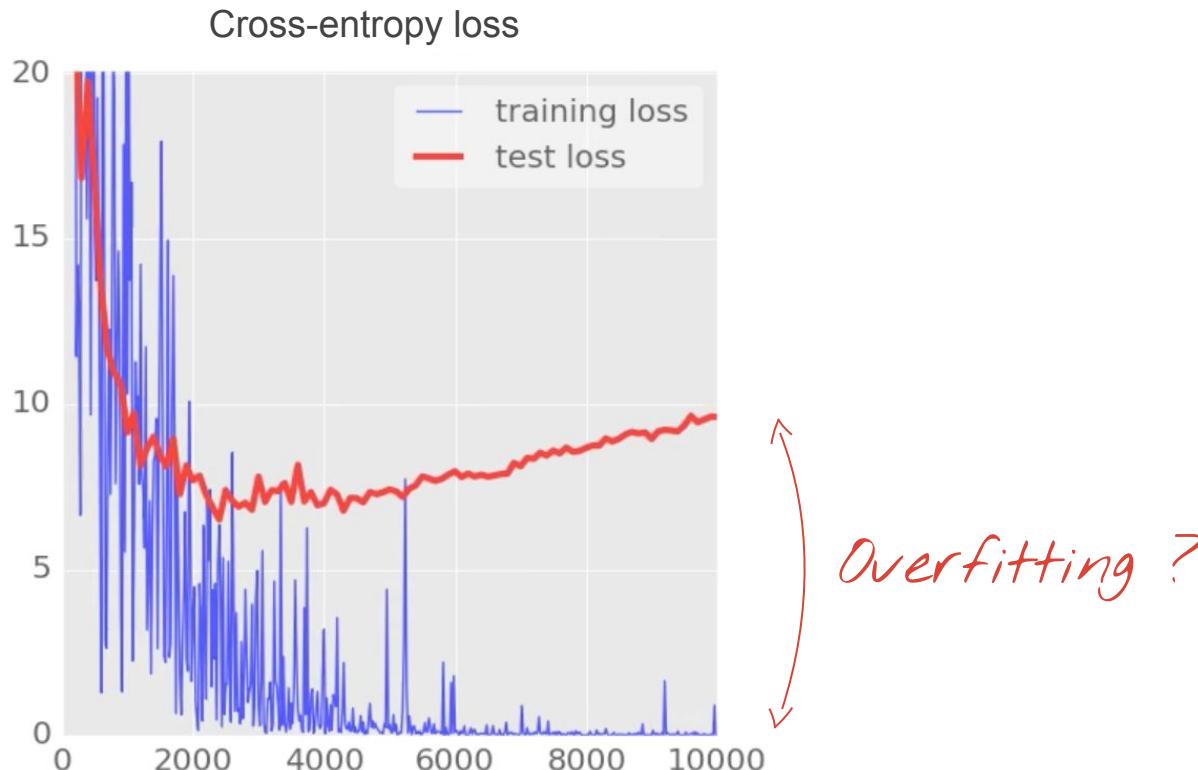


Learning rate decay



Learning rate 0.003 at start then dropping exponentially to 0.0001

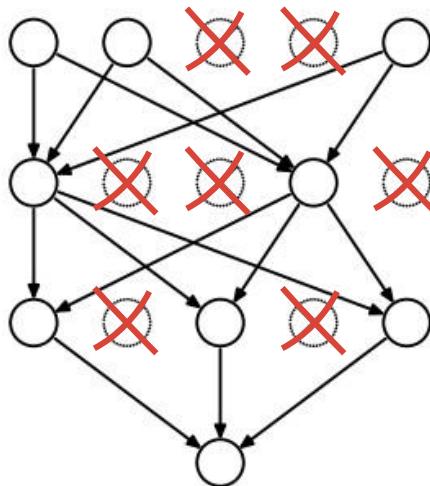
Overfitting



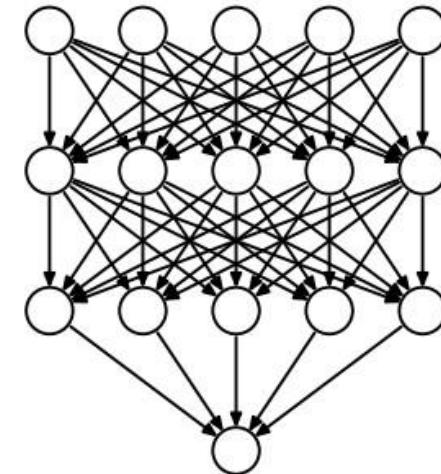
Dropout



Dropout



*TRAINING
rate=0.5*

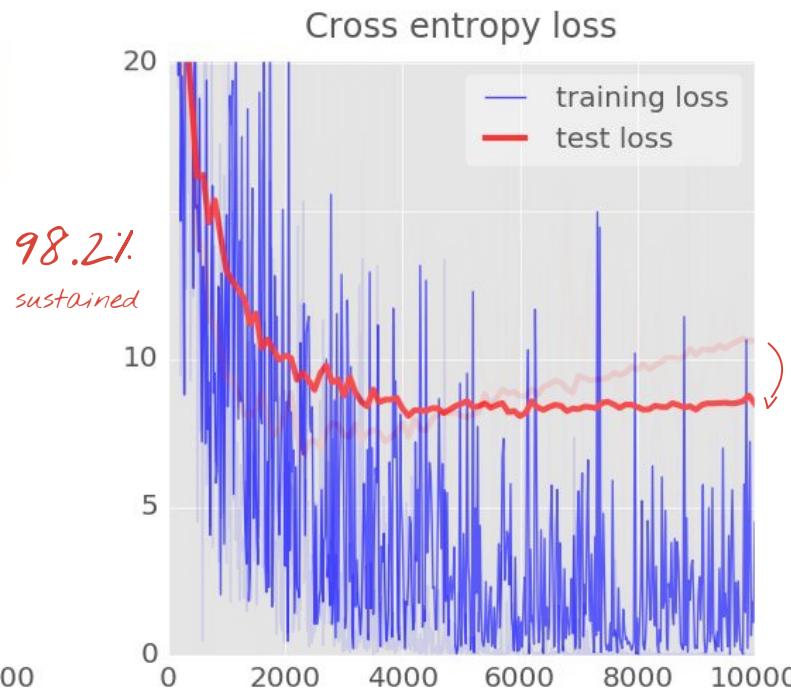
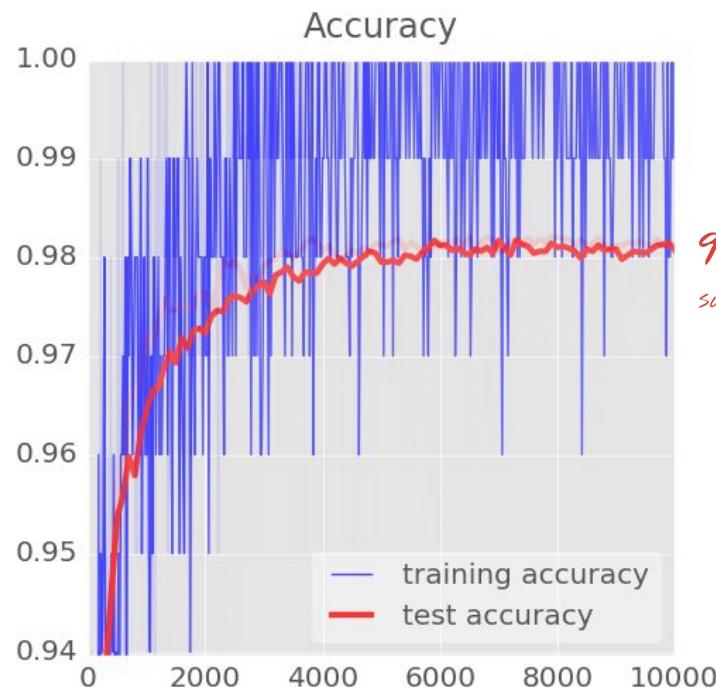


*EVALUATION
rate=0*

```
pkeep =  
tf.placeholder(tf.float32)
```

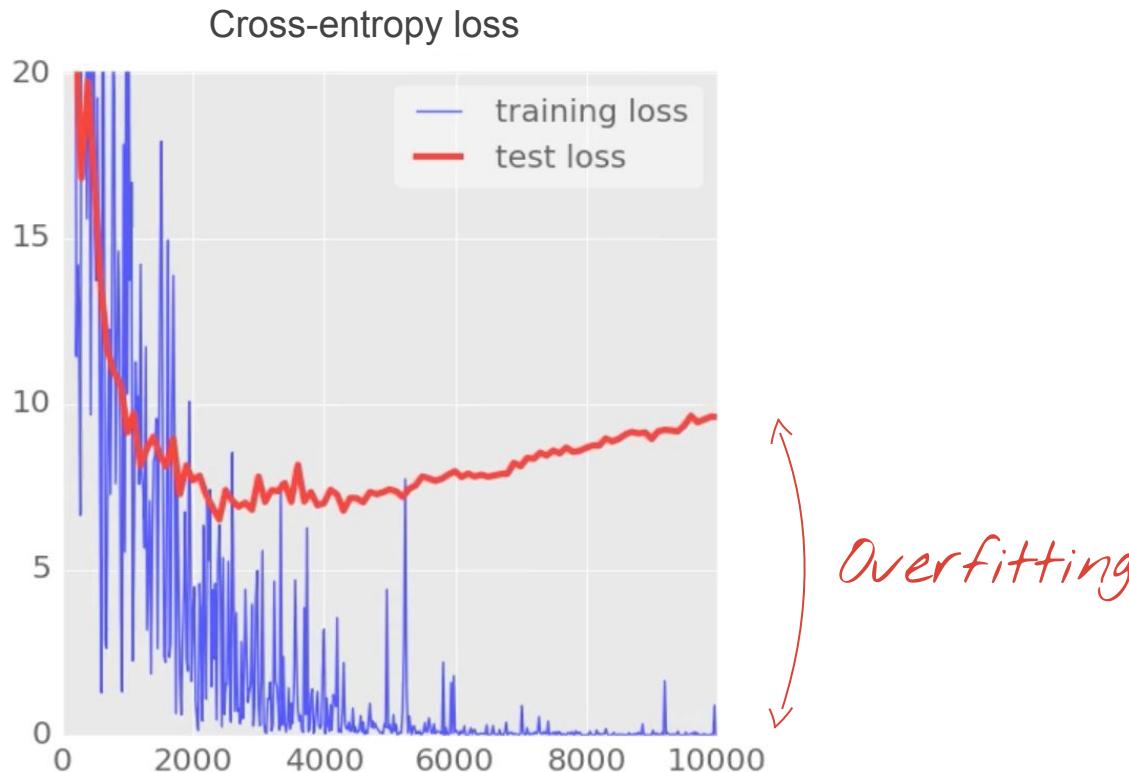
```
Yf = tf.nn.relu(tf.matmul(X, W) + B)  
Y = tf.nn.dropout(Yf, pkeep)
```

All the party tricks



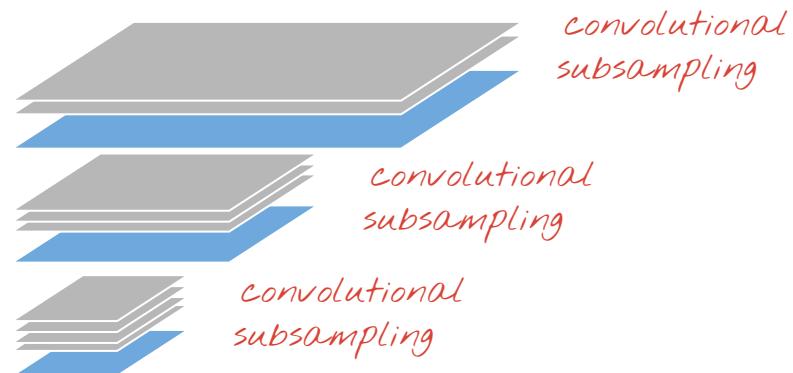
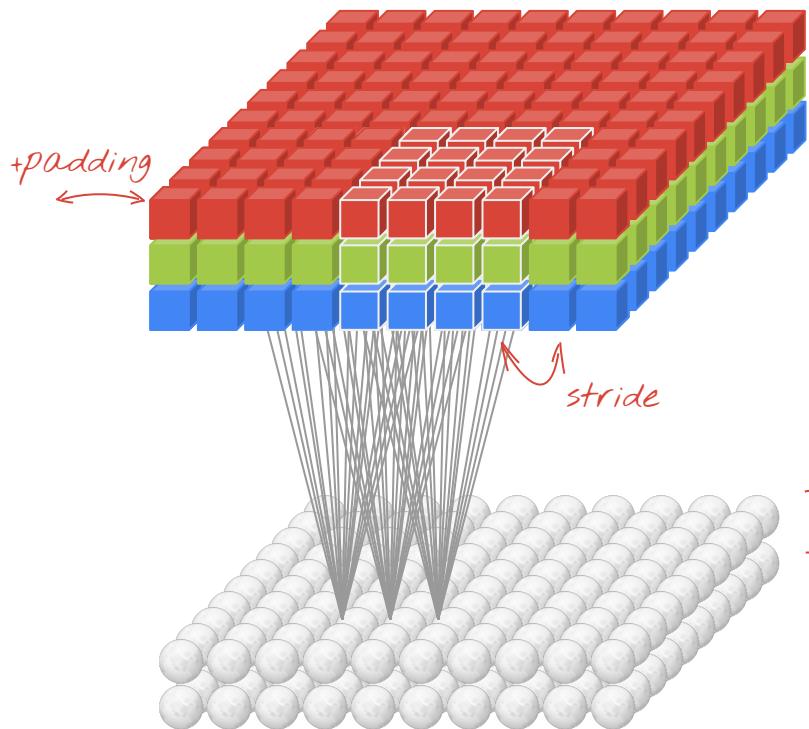
RELU, decaying learning rate 0.003 → 0.0001 and dropout 0.75

Overfitting





Convolutional layer



$W[4, 4, 3]$
 $W_2[4, 4, 3]$ | $W[4, 4, 3, 2]$

filter size input channels output channels

Hacker's tip

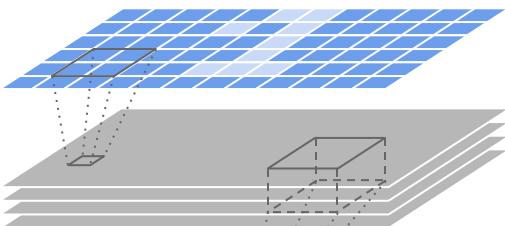


ALL
Convolu-
tional

Convolutional neural network

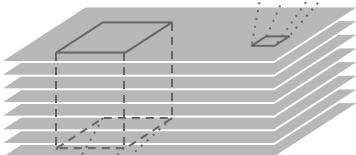
+ biases on
all layers

$28 \times 28 \times 1$

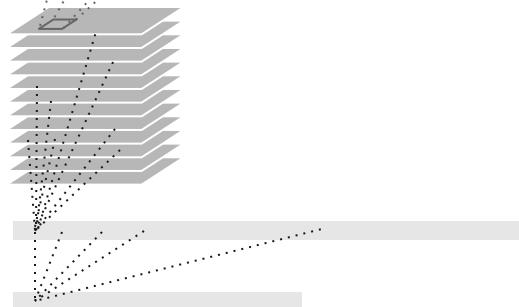


$28 \times 28 \times 4$

$14 \times 14 \times 8$



$7 \times 7 \times 12$



200

10

convolutional layer, 4 channels

$W1[5, 5, 1, 4]$ stride 1

convolutional layer, 8 channels

$W2[4, 4, 4, 8]$ stride 2

convolutional layer, 12 channels

$W3[4, 4, 8, 12]$ stride 2

fully connected layer

$W4[7 \times 7 \times 12, 200]$

softmax readout layer

$W5[200, 10]$

Tensorflow - initialisation

K=4

L=8

M=12

*filter
size* *input
channels* *output
channels*

W1 = tf.Variable(tf.truncated_normal([5, 5, 1, K], stddev=0.1))

B1 = tf.Variable(tf.ones([K])/10)

W2 = tf.Variable(tf.truncated_normal([5, 5, K, L], stddev=0.1))

B2 = tf.Variable(tf.ones([L])/10)

W3 = tf.Variable(tf.truncated_normal([4, 4, L, M], stddev=0.1))

B3 = tf.Variable(tf.ones([M])/10)

N=200

*weights initialised
with random values*

W4 = tf.Variable(tf.truncated_normal([7*7*M, N], stddev=0.1))

B4 = tf.Variable(tf.ones([N])/10)

W5 = tf.Variable(tf.truncated_normal([N, 10], stddev=0.1))

B5 = tf.Variable(tf.zeros([10])/10)

Tensorflow - the model

input image batch
 $X[100, 28, 28, 1]$

weights

stride

biases

```
Y1 = tf.nn.relu(tf.nn.conv2d(X, W1, strides=[1, 1, 1, 1], padding='SAME') + B1)
Y2 = tf.nn.relu(tf.nn.conv2d(Y1, W2, strides=[1, 2, 2, 1], padding='SAME') + B2)
Y3 = tf.nn.relu(tf.nn.conv2d(Y2, W3, strides=[1, 2, 2, 1], padding='SAME') + B3)
```

```
YY = tf.reshape(Y3, shape=[-1, 7 * 7 * M])
```

```
Y4 = tf.nn.relu(tf.matmul(YY, W4) + B4)
```

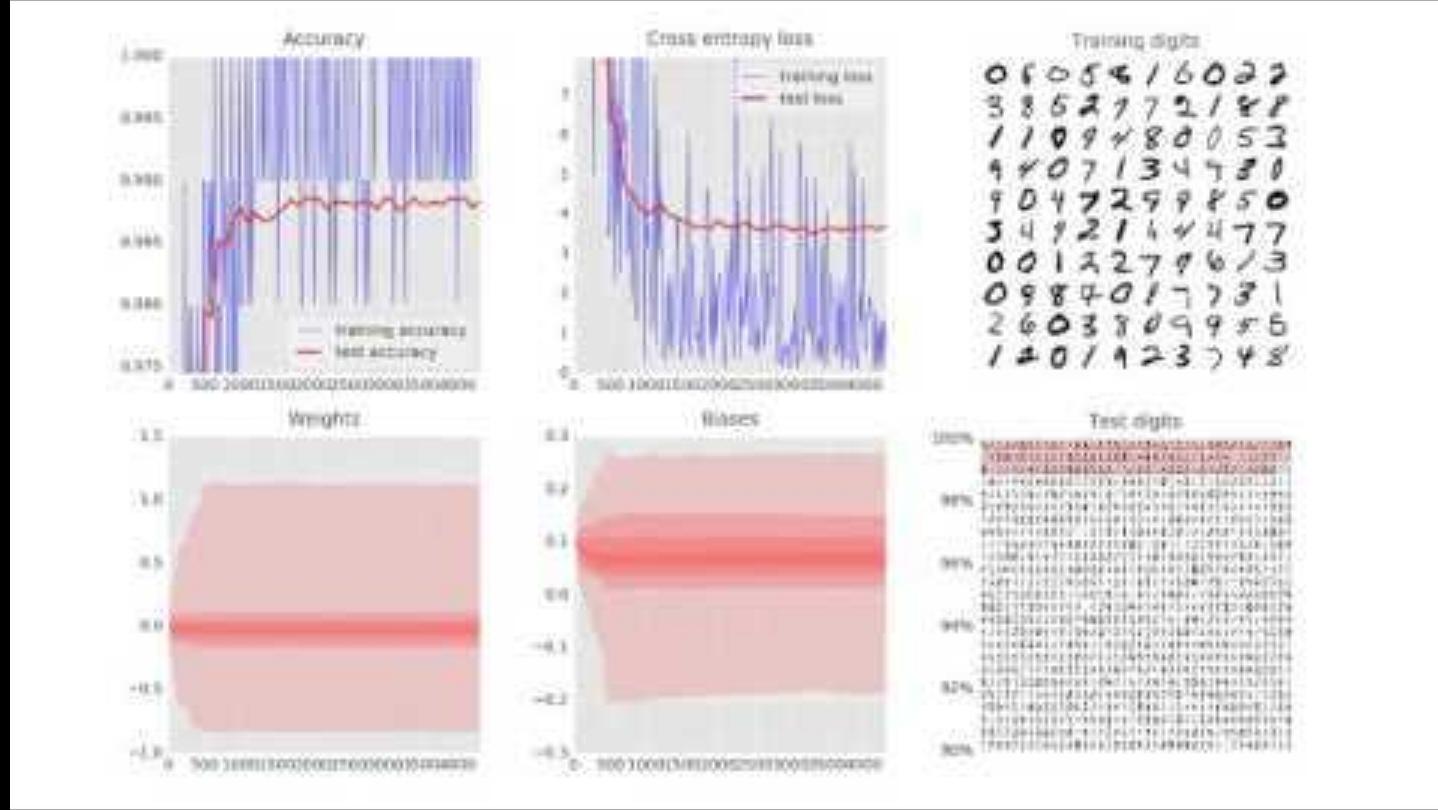
```
Y  = tf.nn.softmax(tf.matmul(Y4, W5) + B5)
```

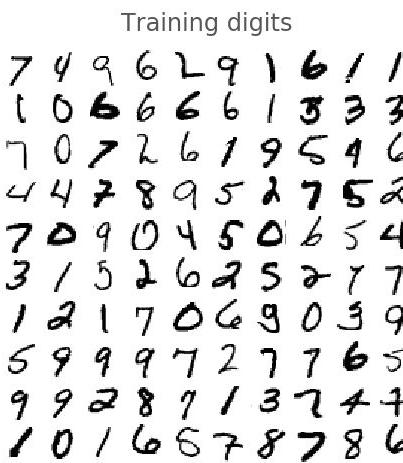
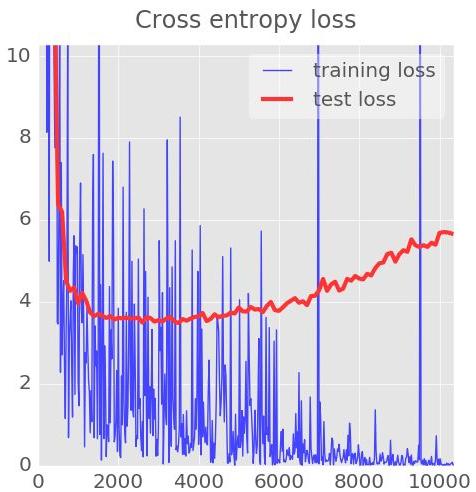
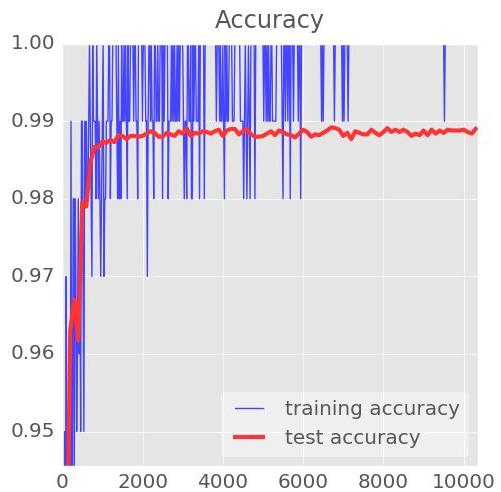
flatten all values for
fully connected layer

$Y3 [100, 7, 7, 12]$

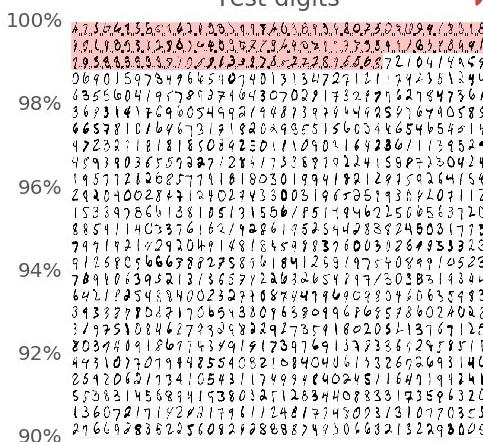
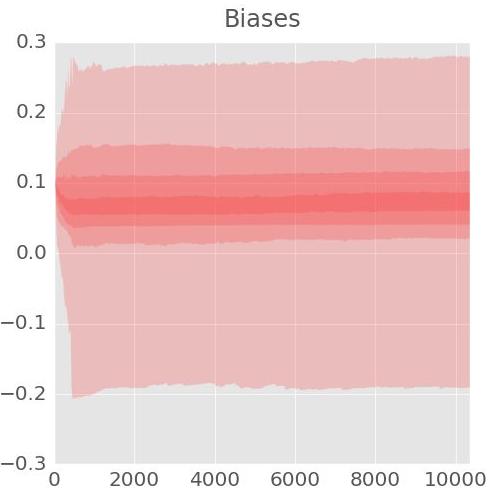
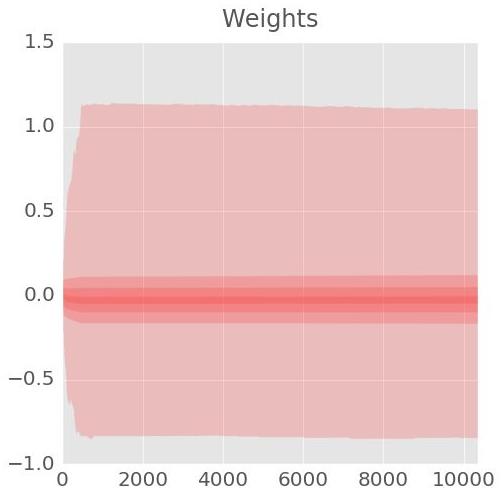
$YY [100, 7x7x12]$

Demo

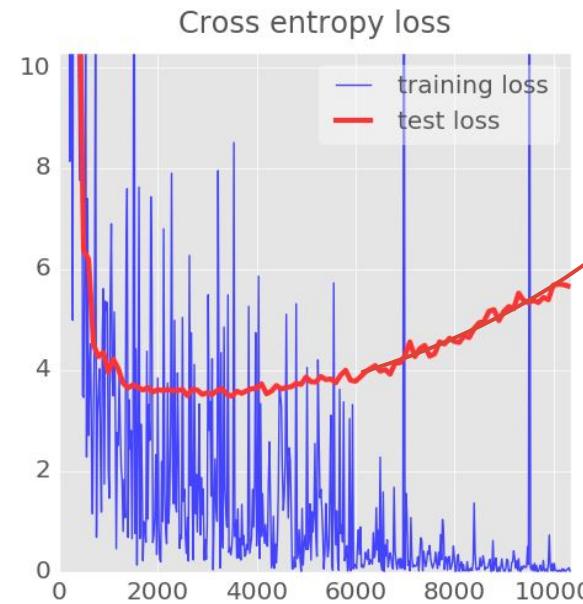
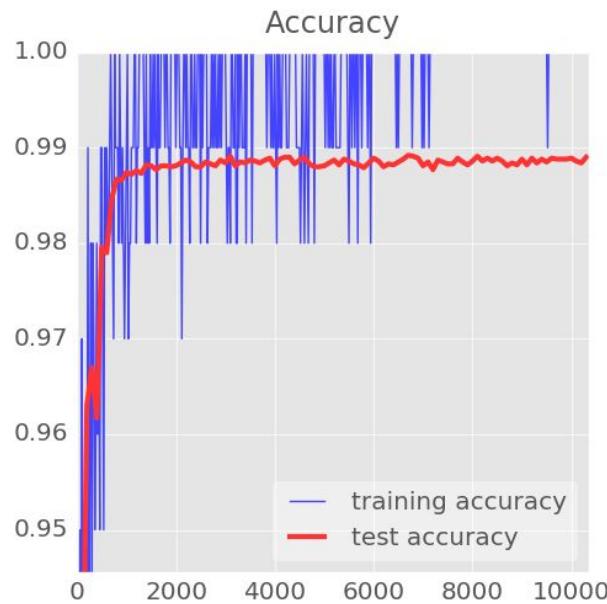




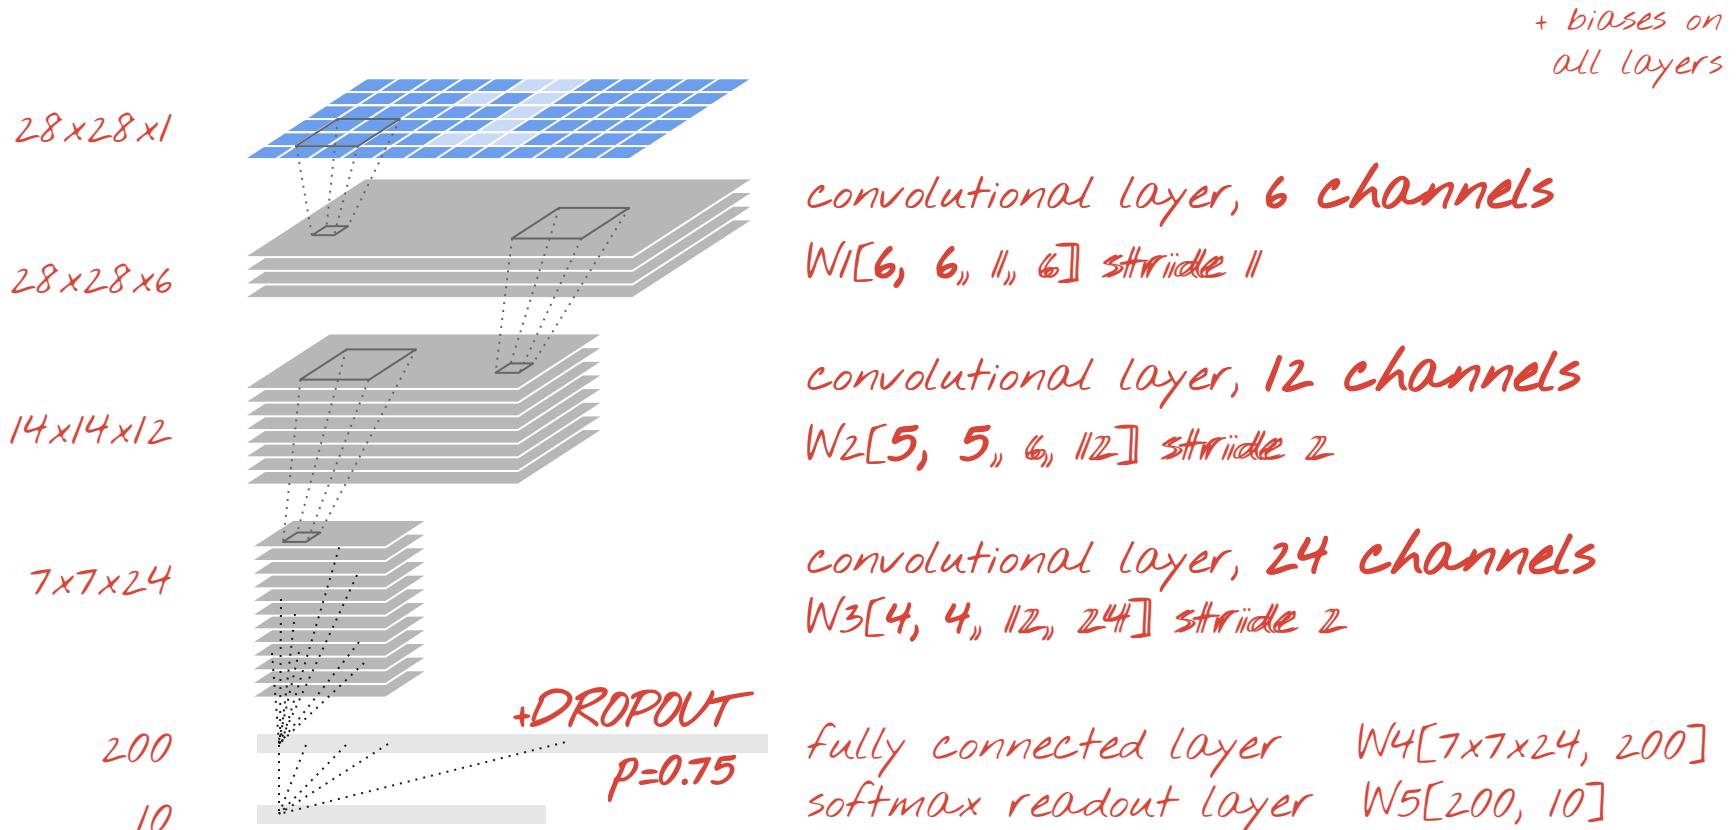
98.9%



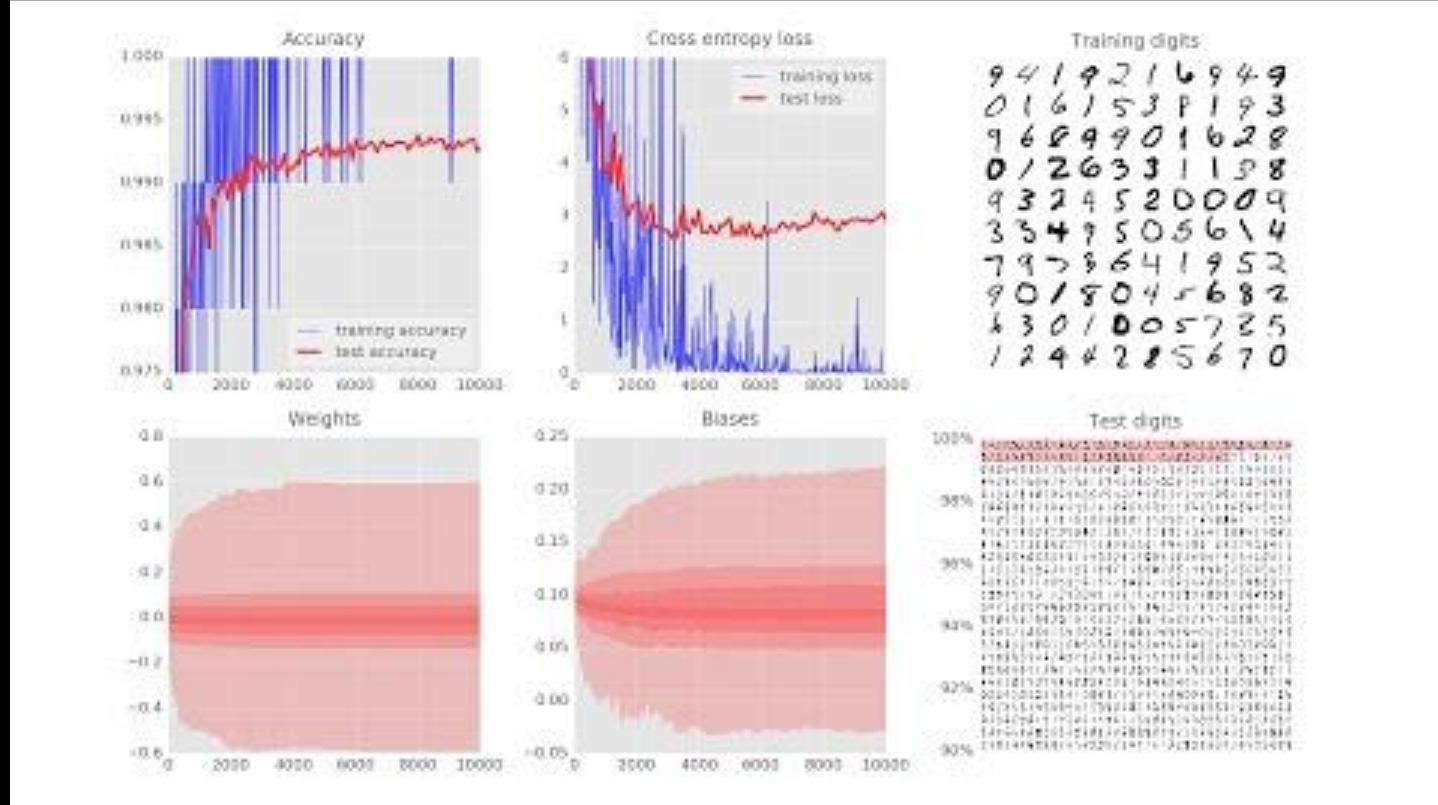
WTXH ???

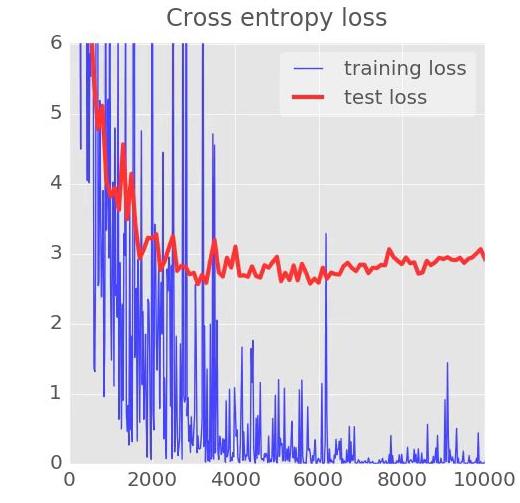
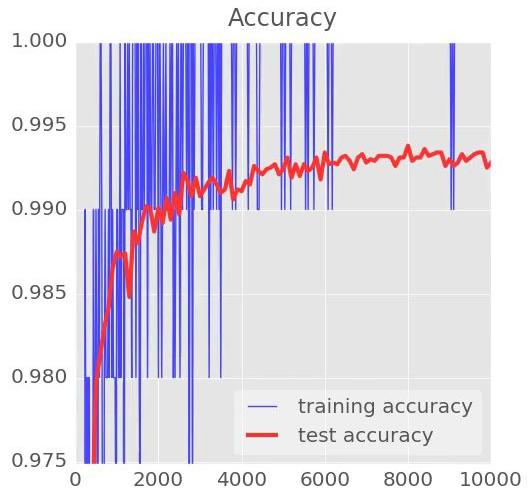


Bigger convolutional network + dropout

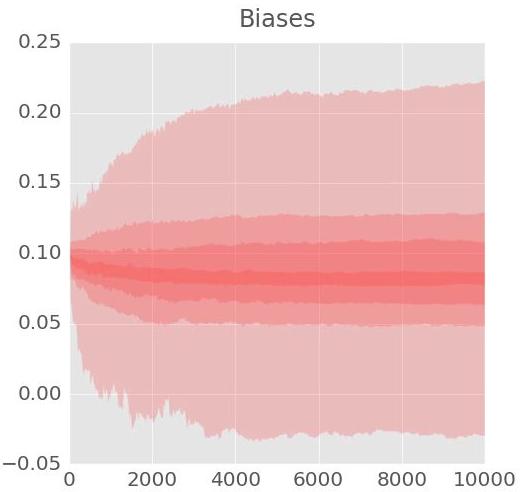
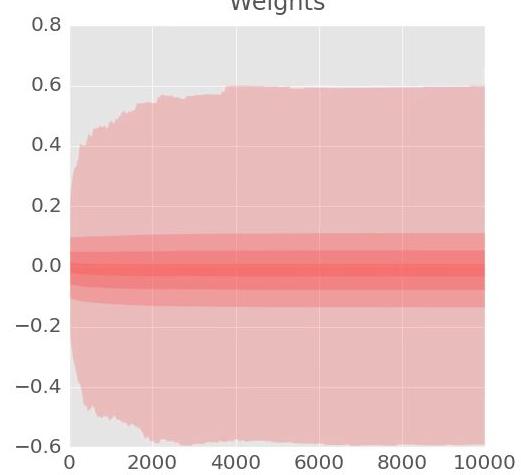


Demo



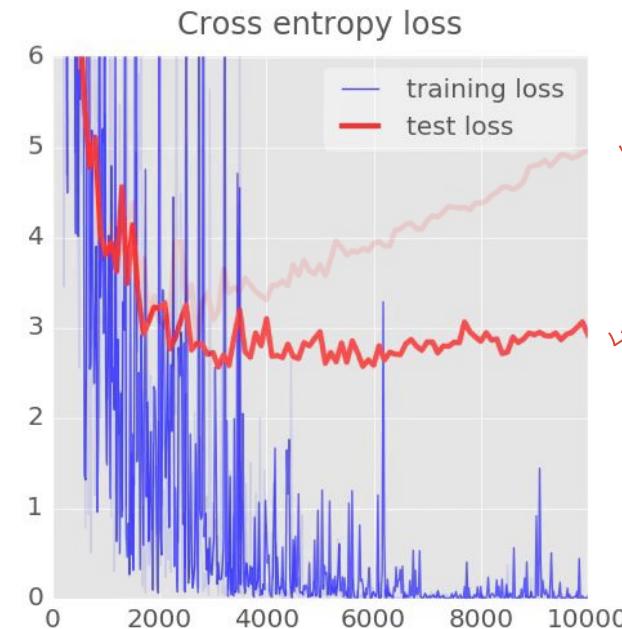
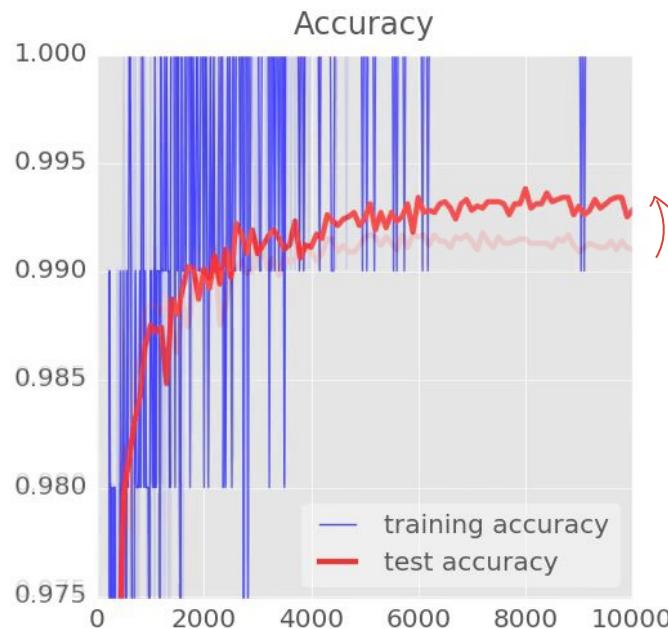


99.3%



100%	6475323215243131037401231143102131003110310 6211321351031021310131031103102131031021310310 6706901597394916490140131347211111943113 446353560419737164307027113217112112112112112 6436293116169605499014813371442501674958 5665331016967317162034955154034465495821 447252716181860842501110931642367113152 44839130365598872841133891922411594735042 419577128268771161805019941922911924115 629200002347134027133003192581131120111 21533928631135105121556755114946225065637 2008541140337616211286119534428362460317 737911927292019148118159913750012649332 36116905663829758361691259197403991052 37849102392131365712481126581171305821133 4642110541340023251687441194901804163348 339337701706543109635099176557860240223 1795118442147938982927359118020511311125 20394601130714517911735176913123361291811 443107019144285405216450061326206314 62512062113410543117493446402451164919124 1513631245689415380381128344688331935632 613607211712411111248119149003181071035 34166128322351092128086874931063211229560 5131416029117473938347121223233991190358
98%	6475323215243131037401231143102131003110310 6211321351031021310131031103102131031021310310 6706901597394916490140131347211111943113 446353560419737164307027113217112112112112112 6436293116169605499014813371442501674958 5665331016967317162034955154034465495821 447252716181860842501110931642367113152 44839130365598872841133891922411594735042 419577128268771161805019941922911924115 629200002347134027133003192581131120111 21533928631135105121556755114946225065637 2008541140337616211286119534428362460317 737911927292019148118159913750012649332 36116905663829758361691259197403991052 37849102392131365712481126581171305821133 4642110541340023251687441194901804163348 339337701706543109635099176557860240223 1795118442147938982927359118020511311125 20394601130714517911735176913123361291811 443107019144285405216450061326206314 62512062113410543117493446402451164919124 1513631245689415380381128344688331935632 613607211712411111248119149003181071035 34166128322351092128086874931063211229560 5131416029117473938347121223233991190358
96%	6475323215243131037401231143102131003110310 6211321351031021310131031103102131031021310310 6706901597394916490140131347211111943113 446353560419737164307027113217112112112112112 6436293116169605499014813371442501674958 5665331016967317162034955154034465495821 447252716181860842501110931642367113152 44839130365598872841133891922411594735042 419577128268771161805019941922911924115 629200002347134027133003192581131120111 21533928631135105121556755114946225065637 2008541140337616211286119534428362460317 737911927292019148118159913750012649332 36116905663829758361691259197403991052 37849102392131365712481126581171305821133 4642110541340023251687441194901804163348 339337701706543109635099176557860240223 1795118442147938982927359118020511311125 20394601130714517911735176913123361291811 443107019144285405216450061326206314 62512062113410543117493446402451164919124 1513631245689415380381128344688331935632 613607211712411111248119149003181071035 34166128322351092128086874931063211229560 5131416029117473938347121223233991190358
94%	6475323215243131037401231143102131003110310 6211321351031021310131031103102131031021310310 6706901597394916490140131347211111943113 446353560419737164307027113217112112112112112 6436293116169605499014813371442501674958 5665331016967317162034955154034465495821 447252716181860842501110931642367113152 44839130365598872841133891922411594735042 419577128268771161805019941922911924115 629200002347134027133003192581131120111 21533928631135105121556755114946225065637 2008541140337616211286119534428362460317 737911927292019148118159913750012649332 36116905663829758361691259197403991052 37849102392131365712481126581171305821133 4642110541340023251687441194901804163348 339337701706543109635099176557860240223 1795118442147938982927359118020511311125 20394601130714517911735176913123361291811 443107019144285405216450061326206314 62512062113410543117493446402451164919124 1513631245689415380381128344688331935632 613607211712411111248119149003181071035 34166128322351092128086874931063211229560 5131416029117473938347121223233991190358
92%	6475323215243131037401231143102131003110310 6211321351031021310131031103102131031021310310 6706901597394916490140131347211111943113 446353560419737164307027113217112112112112112 6436293116169605499014813371442501674958 5665331016967317162034955154034465495821 447252716181860842501110931642367113152 44839130365598872841133891922411594735042 419577128268771161805019941922911924115 629200002347134027133003192581131120111 21533928631135105121556755114946225065637 2008541140337616211286119534428362460317 737911927292019148118159913750012649332 36116905663829758361691259197403991052 37849102392131365712481126581171305821133 4642110541340023251687441194901804163348 339337701706543109635099176557860240223 1795118442147938982927359118020511311125 20394601130714517911735176913123361291811 443107019144285405216450061326206314 62512062113410543117493446402451164919124 1513631245689415380381128344688331935632 613607211712411111248119149003181071035 34166128322351092128086874931063211229560 5131416029117473938347121223233991190358
90%	6475323215243131037401231143102131003110310 6211321351031021310131031103102131031021310310 6706901597394916490140131347211111943113 446353560419737164307027113217112112112112112 6436293116169605499014813371442501674958 5665331016967317162034955154034465495821 447252716181860842501110931642367113152 44839130365598872841133891922411594735042 419577128268771161805019941922911924115 629200002347134027133003192581131120111 21533928631135105121556755114946225065637 2008541140337616211286119534428362460317 737911927292019148118159913750012649332 36116905663829758361691259197403991052 37849102392131365712481126581171305821133 4642110541340023251687441194901804163348 339337701706543109635099176557860240223 1795118442147938982927359118020511311125 20394601130714517911735176913123361291811 443107019144285405216450061326206314 62512062113410543117493446402451164919124 1513631245689415380381128344688331935632 613607211712411111248119149003181071035 34166128322351092128086874931063211229560 5131416029117473938347121223233991190358

YEAH !



with dropout



Have fun !



Martin Görner

Google Developer relations
[@martin_gorner](https://twitter.com/martin_gorner)



TensorFlow

Videos, slides, code:

[github.com/
GoogleCloudPlatform/
tensorflow-without-a-phd](https://github.com/GoogleCloudPlatform/tensorflow-without-a-phd)



Cloud ML Engine

your TensorFlow models
trained in Google's cloud.



Cloud Auto ML Vision ALPHA

Just bring your data



Google Cloud



Cloud TPU

ML supercomputing

Pre-trained models:



Cloud Vision API



Cloud Speech API



Natural Language API



Google Translate API



Video Intelligence API



Cloud Jobs API BETA



>TensorFlow and deep learning_ without a PhD



>TensorFlow, deep learning and \\ recurrent neural networks without a PhD



>TensorFlow, deep learning and \\ modern convolutional neural nets without a PhD



>TensorFlow and \\ deep reinforcement learning without a PhD



The superpower: batch normalisation



>TensorFlow, deep learning and \\ modern RNN architectures without a PhD



Tensorflow and
deep learning
without a PhD

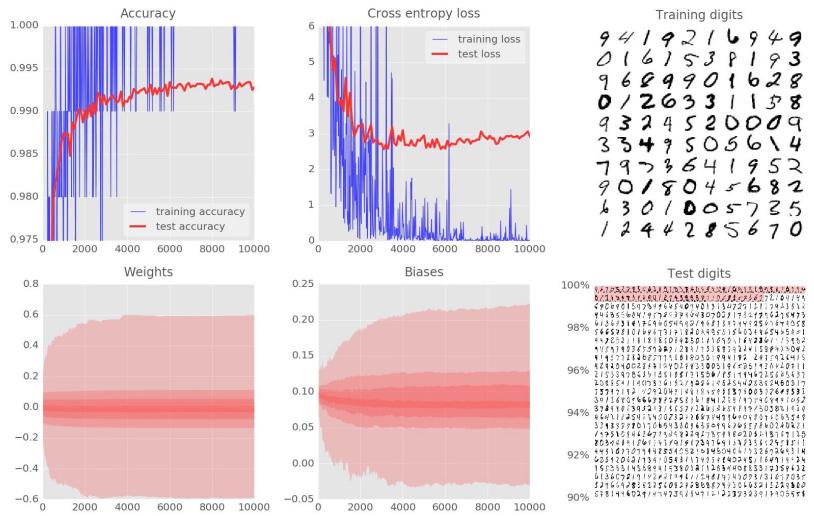


@martin_gorner

github.com/
GoogleCloudPlatform/
tensorflow-without-a-phd



Workshop



If training speed is an issue (it can happen in VirtualBox), consider displaying graph 3 while you wait instead of all the six graphs. You can also disable the visualisations altogether. There are instructions for that purpose at the end of each code sample.

Keyboard shortcuts for the visualisation GUI:

1	display 1 st graph only
2	display 2 nd graph only
3	display 3 rd graph only
4	display 4 th graph only
5	display 5 th graph only
6	display 6 th graph only
7	display graphs 1 and 2
8	display graphs 4 and 5
9	display graphs 3 and 6
ESC or 0 ..	back to displaying all graphs
SPACE	pause/resume
0	box zoom mode (then use mouse)
H	reset all zooms
Ctrl-S	save current image

Workshop

Self-paced code lab (summary below ↓): goo.gl/mVZloU
Code: github.com/martin-gorner/tensorflow-mnist-tutorial



I-5. Theory ([install](#) then sit back and listen or [read](#))

Neural networks 101: softmax, cross-entropy, mini-batching, gradient descent, hidden layers, sigmoids, and how to implement them in Tensorflow



6. Practice ([full instructions for this step](#))

Open file: `mnist_1.0_softmax.py`
Run it, play with the visualisations (keyboard shortcuts on previous slide), read and understand the code as well as the basic structure of a Tensorflow program.



7. Practice ([full instructions for this step](#))

Start from the file `mnist_1.0_softmax.py` and add one or two hidden layers.

Solution in: `mnist_2.0_five_layers_sigmoid.py`



8. Practice ([full instructions for this step](#))

Special care for deep neural networks: use RELU activation functions, use a better optimiser, initialise weights with random values and beware of the log(0)



9-10. Practice ([full instructions for this step](#))

Use a decaying learning rate and then add dropout

Solution in: `mnist_2.2_five_layers_relu_lrdecay_dropout.py`



II. Theory (sit back and listen or [read](#))

Convolutional networks



12. Practice ([full instructions for this step](#))

Replace your model with a convolutional network, without dropout.

Solution in: `mnist_3.0_convolutional.py`



13. Challenge ([full instructions for this step](#))

Try a bigger neural network (good hyperparameters on slide 43) and add dropout on the last layer to get >99%

Solution in: `mnist_3.0_convolutional_bigger_dropout.py`

